

Introduction to Database Systems

2023-Fall

What is needed for query optimization?

- **Given: A closed set of operators**
 - Relational ops (table in, table out)
 - Physical implementations (of those ops and a few more)

1. Plan space

- Based on relational equivalences, different implementations

2. Cost Estimation based on

- Cost formulas
- Size estimation, in turn based on
 - Catalog information on base tables
 - Selectivity (Reduction Factor) estimation
- Target: To sift through the plan space and find lowest cost option!

Plan Space Review

- For a SQL query, full plan space:
 - All equivalent relational algebra expressions
 - Based on the equivalence rules we learned
 - All mixes of physical implementations of those algebra expressions
- We might prune this space:
 - Selection/Projection pushdown
 - Avoid cartesian products

Plan Space: Too Large, must be pruned.

- Only the space of *left-deep plans* is considered.
- Left-deep plans allow output of each operator to be *pipelined* into the next operator without storing it in a temporary relation.
- Cartesian products avoided.

Schema for Examples

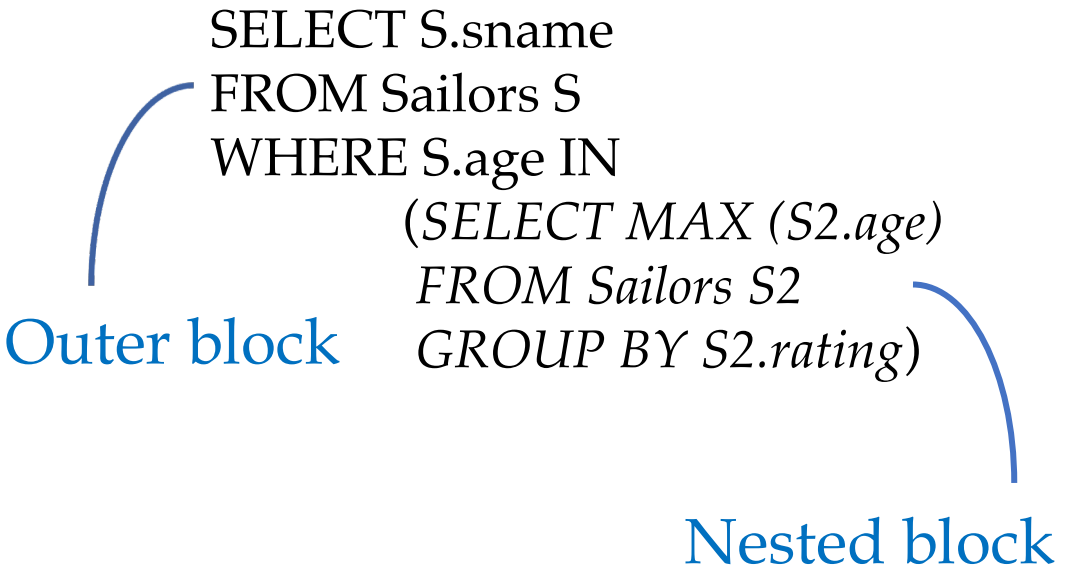
Sailors (*sid*: integer, *sname*: text, *rating*: integer, *age*: float)

Reserves (*sid*: integer, *bid*: integer, *day*: date, *rname*: text)

- Reserves:
 - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
 - 100 distinct bids.
- Sailors:
 - Each tuple is 50 bytes long,
 - 80 tuples per page, 500 pages.
 - 10 ratings, 40,000 sids.

Query Blocks: Units of Optimization

- An SQL query is parsed into a collection of **query blocks**, and these are optimized one block at a time.
- Nested blocks are usually treated as calls to a subroutine, made once per outer tuple. (This is an over-simplification, but serves for now.)
- For each block, the plans considered are:
 - All *left-deep join trees* (all ways to join the relations one-at-a-time)



SELECT S.sname
FROM Sailors S
WHERE S.age IN
 (SELECT MAX (S2.age)
 FROM Sailors S2
 GROUP BY S2.rating)

Outer block

Nested block

The diagram shows an SQL query with two nested blocks highlighted by blue arcs. The outer block is the entire query, and the nested block is the subquery in the WHERE clause. The labels 'Outer block' and 'Nested block' are placed below the arcs.

Cost Estimation

- For each plan considered, must estimate total cost:
 - Must estimate **cost** of each operation in plan tree.
 - Depends on input cardinalities.
 - We've already discussed this for various operators
 - sequential scan, index scan, joins, etc.
 - Must estimate **size of result** for each operation in tree!
 - Because it determines downstream input cardinalities!
 - Use information about the input relations.
 - For selections and joins, assume independence of predicates.

Statistics and Catalogs

- Need info on relations and indexes involved.

- **Catalogs** typically contain at least:

| Statistic | Meaning |
|-----------|--------------------------------------|
| NTuples | # of tuples in a table (cardinality) |
| NPages | # of disk pages in a table |
| Low/High | min/max value in a column |
| Nkeys | # of distinct values in a column |
| IHeight | the height of an index |
| INPages | # of disk pages in an index |


- Catalogs updated periodically.
 - Too expensive to do continuously
 - Lots of approximation anyway, so a little slop here is ok.
- Modern systems do more
 - Esp. keep more detailed statistical information on data values
 - e.g., *histograms*

Cost Estimates for Single-Relation Plans

- Index I on primary key matches selection:
 - *Cost is $\text{Height}(I)+1$ for a B+ tree, about 1.2 for hash index.*
- Clustered index I matching one or more selects:
 - *$(\text{NPages}(I)+\text{NPages}(R)) * \text{product of RF's of matching selects.}$*
- Non-clustered index I matching one or more selects:
 - *$(\text{NPages}(I)+\text{NTuples}(R)) * \text{product of RF's of matching selects.}$*
- Sequential scan of file:
 - *$\text{NPages}(R)$.*
- + **Note:** *Typically, no duplicate elimination on projections! (Exception: Done on answers if user says DISTINCT.)*

Queries Over Multiple Relations

- Fundamental decision in System R: *only left-deep join trees* are considered.
- As the number of joins increases, the number of alternative plans grows rapidly; ***we need to restrict the search space***. Left-deep trees allow us to generate all *fully pipelined* plans.
- Intermediate results not written to temporary files.
Not all left-deep trees are fully pipelined (e.g., SortMerge join).



Enumeration of Left-Deep Plans

- Left-deep plans differ only in the order of relations, the access method for each relation, and the join method for each join.
- Enumerated using N passes (if N relations joined):
 - Pass 1: Find best 1-relation plan for each relation.
 - Pass 2: Find best way to join result of each 1-relation plan (as outer) to another relation. (*All 2-relation plans.*)
 - Pass N: Find best way to join result of a (N-1)-relation plan (as outer) to the N'th relation. (*All N-relation plans.*)
- For each subset of relations, retain only:
 - Cheapest plan overall, plus
 - Cheapest plan for each *interesting order* of the tuples.

Enumeration of Plans (Contd.)

- ORDER BY, GROUP BY, aggregates etc. handled as a final step, using either an 'interestingly ordered' plan or an additional sorting operator.
- An N-1 way plan is not combined with an additional relation unless there is a join condition between them, unless all predicates in WHERE have been used up.
 - i.e., avoid Cartesian products if possible.
- In spite of pruning plan space, this approach is still exponential in the # of tables.

Cost Estimation for Multi-relation Plans

- Consider a query block:
- Maximum # tuples in result is the product of the cardinalities of relations in the FROM clause.
- ***Reduction factor (RF)*** associated with each *term* reflects the impact of the *term* in reducing result size. *Result cardinality* = Max # tuples * product of all RF's.
- Multi-relation plans are built up by joining one new relation at a time.
 - Cost of join method, plus estimation of join cardinality gives us both cost estimate and result size estimate